

Lightning-Fast IPC through Lock-Free Shared Memory

Daniel Fox Franke <dfoxfranke@gmail.com>

Purveyors of Fine Gadgets

June 2024

Intro quiz

Global

```
int x = 1;
```

Thread 1

```
x = 2;
```

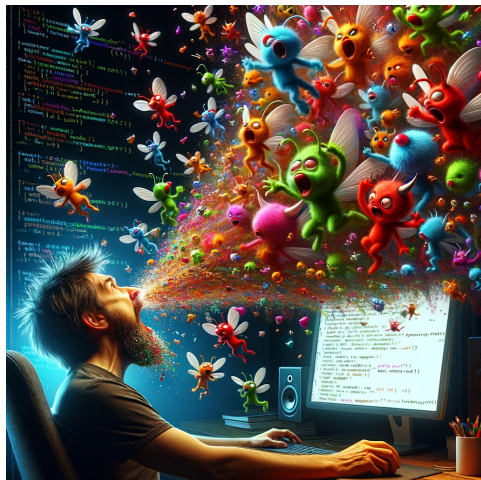
Thread 2

```
int a, b;  
a = x;  
b = x;  
printf("%d %d", a, b);
```

What can this print?

- 1 1?
- 1 2?
- 2 2?
- 2 1?
- Something else?

Answer: all of the above



C17 §5.1.2.4.35:

“The execution of a program contains a *data race* if it contains two **conflicting actions** in different threads, at least one of which is not **atomic**, and neither **happens before** the other. Any such data race results in undefined behavior.”

Avoiding UB with atomics

Global

```
#include <stdatomic.h>
atomic_int x;
atomic_init(&x, 1);
```

Thread 1

```
atomic_store(&x, 2);
```

Thread 2

```
int a, b;
a = atomic_load(&x);
b = atomic_load(&x);
printf("%d %d", a, b);
```

What can this print?

✓ 1 1

✓ 1 2

✓ 2 2

✗ 2 1

✗ Something else

Same thing in Rust

Global

```
use std::sync::atomic::{AtomicI32, Ordering::SeqCst};  
static const x : AtomicI32 = AtomicI32::new(1);
```

Thread 1

```
x.store(2, SeqCst);
```

Thread 2

```
let a = x.load(SeqCst);  
let b = x.load(SeqCst);  
print!("{a} {b}");
```

What can this print?

✓ 1 1


✓ 1 2

✓ 2 2

✗ 2 1

✗ Something else

Cache hierarchy



The ideal: write-back LRU cache

On read:

- 1 Is the word already in cache?
Return it and done.
- 2 Otherwise, store least recently used word to parent cache.
- 3 Load from parent cache into space just freed up, then return it.

On write:

- 1 Is the word already in cache?
Overwrite it and done.
- 2 Otherwise, store least recently used word to parent cache.
- 3 Write to space just freed up. Do *not* immediately forward this new value to parent.

The reality: N -way associative cache

Big LRUs are expensive to implement in hardware. It's cheaper (lower gate count, more energy-efficient) to have lots of little ones. This leads to some loss of performance, but usually negligible.

Example: 4 MiB 8-way associative cache w/ 64-byte cache lines.

- 4 MiB = 2^{16} 64-byte lines.
- Cache has 2^{13} LRUs that hold 8 lines each.
- Only one LRU can be used for any given address. Bits 18:6 of address determine which one.

Cache coherence: problem statement

Write-back caching + multiple cores w/private cache + shared memory = chaos

To keep our sanity, we need:

- 1 **Change propagation:** Writes by one core become visible to all other cores in a *reasonably* timely manner.
- 2 **Sequential consistency:** Writes to any *single* memory location happen in a globally consistent order.

Single-location sequential consistency

Global

```
#define store_relaxed(p, v) atomic_store_explicit( \
    p, v, memory_order_relaxed);
#define load_relaxed(p) atomic_load_explicit( \
    p, memory_order_relaxed);
atomic_int x;
atomic_init(&x, 0);
```

Thread 1

```
store_relaxed(&x, 1);
```

Thread 2

```
store_relaxed(&x, 2);
```

Threads 3 & 4

```
int a, b;
a = load_relaxed(&x);
b = load_relaxed(&x);
printf("%d %d", a, b);
```

Could print:

✓ any of seven possibilities, including...

✓ **1 2**

✓ **2 1**

But **never these two at the same time!**

Sequential *in*consistency with multiple addresses

Global

```
#define store_relaxed(p, v) atomic_store_explicit( \
    p, v, memory_order_relaxed);
#define load_relaxed(p) atomic_load_explicit( \
    p, memory_order_relaxed);
atomic_int x, y;
atomic_init(&x, 1);
atomic_init(&y, 1);
```

Thread 1

```
store_relaxed(&x, 2);
store_relaxed(&y, 2);
```

Threads 2 & 3

```
int b = load_relaxed(&y);
int a = load_relaxed(&x);
printf("%d %d", a, b);
```

Could print:

✓ 1 1

✓ 2 2

✓ **1 2**

✓ **2 1**

Simultaneously!

Cache coherence protocols

- Simplest possible solution: write-through instead of write-back
- More bandwidth-efficient: MESI protocol. Snoop on other cores' memory accesses and maintain state per cache line:
 - **Modified** - Cache line dirty (a write is pending). All other cores in invalid state for this address.
 - **Exclusive** - Cache line clean. All other cores in invalid state for this address.
 - **Shared** - Cache line clean. Other cores may also be in shared state for this address.
 - **Invalid** - Cache line treated as empty; must load from parent before use.

State machine transitions happen on issuance of load/store instructions, or on observation of bus activity.

Exercise: draw this state machine on your own.

CPU support for concurrency: fences

Most architectures have some equivalent to these instructions:

- **Load fence:** all **loads** before the fence instruction execute before any **loads/stores** after it.
- **Store fence:** all **loads/stores** before the fence instruction execute before any **stores** after it.
- **Full fence:** all **loads/stores** before the fence instruction execute before any **loads/stores** after it.

A full fence is not just two half fences! Only a full fence prevents earlier stores from being reordered with later loads.

CPU support for concurrency: load-acquire/store-release

- **load-acquire:** A load that won't be reordered past later instructions.
- **store-release:** A store that won't be reordered past earlier instructions.

These are weaker (and therefore more efficient) than fences: they only affect their own ordering and not that of earlier loads or later stores.

CPU support for concurrency: atomic compare-and-set

Most architectures have an instruction that does this, atomically:

```
bool atomic_cas(intptr_t* addr, intptr_t old, intptr_t new) {  
    if (*addr == old) {  
        *addr = new;  
        return true;  
    } else {  
        return false;  
    }  
}
```

From atomic CAS to atomic anything

```
intptr_t atomically(intptr_t *addr, intptr_t (*f)(intptr_t)) {  
    do {  
        intptr_t old = atomic_load(addr);  
        intptr_t new = f(old);  
    } while (!atomic_cas(addr, old, new));  
    return new;  
}
```


Very rough latency intuition

- Register \leftrightarrow register: 0.25ns
- Register \leftrightarrow private cache: 1–4ns
- Private cache \rightarrow private cache: 6ns
- Register \leftrightarrow shared cache: 25ns
- Register \leftrightarrow main memory: 100ns

Acquire/release and fences don't add communication, but they add interdependency between communications. Later instructions can't start until their dependencies have finished.

The C++ Memory Model

- Abstracts away architecture-specific details about atomicity and ordering.
- Originally for defining `std::atomic` in C++11.
- Backported to C (since C11) and D.
- Adopted by Rust, Zig, probably others.

Combining and reordering non-atomic operations

Source

```
void foo(int *restrict a, int *restrict b) {  
    *a += 1;    *b += 1;    *a += 1; //Tight slide, excuse the formatting  
}
```

x86_64

```
inc    dword ptr [rsi]  
mov    eax, dword ptr [rdi]  
add    eax, 2  
mov    dword ptr [rdi], eax
```

aarch64

```
ldr    w8, [x1]  
ldr    w9, [x0]  
add    w8, w8, #1  
add    w9, w9, #2  
str    w8, [x1]  
str    w9, [x0]
```

riscv64

```
lw      a2, 0(a1)  
lw      a3, 0(a0)  
addiw   a2, a2, 1  
sw      a2, 0(a1)  
addiw   a3, a3, 2  
sw      a3, 0(a0)
```

Atomic operations and memory orders

Types of atomic operation:

- Load (L)
- Store (S)
- Atomic Read-Modify-Write (RMW)

RMW = atomically mutate a memory location, *e.g.*, atomic increment. Can be built on CAS.

Memory orders and (applicability):

- Relaxed (L, S, RMW)
- Acquire (L, RMW)
- Release (S, RMW)
- Acq+Rel (RMW)
- SeqCst (L, S, RMW)

Memory orderings: basic overview

- **Relaxed**: Still atomic, but no memory ordering guarantees beyond cache coherence. Generates ordinary load/store instructions.
- **Acquire**: Generates load-acquire instructions instead of regular loads, therefore won't be reordered with later instructions.
- **Release**: Generates store-release instructions instead of regular stores, therefore won't be reordered with earlier instructions.
- **AcqRel**: combines acquire semantics with release semantics when implementing RMWs.
- **SeqCst**: Makes the operation sequentially consistent with all other SeqCst operations (only), even when multiple memory locations are involved. Usually implemented using full fences.

To keep these slides readable...

`<operation>_<order>(...)`

is a macro for

```
atomic_<operation>_explicit(  
    ...,  
    memory_order_<order>  
);
```

For everything defined in `<stdatomic.h>`.

Sequential *in*consistency with multiple addresses

Global

```
atomic_int x, y;  
atomic_init(&x, 1);  
atomic_init(&y, 1);
```

Thread 1

```
store_relaxed(&x, 2);  
store_relaxed(&y, 2);
```

Threads 2 & 3

```
int b = load_relaxed(&y);  
int a = load_relaxed(&x);  
printf("%d %d", a, b);
```

Could print:

✓ 1 1

✓ 2 2

✓ **1 2**

✓ **2 1**

Simultaneously!

Sequential consistency with multiple addresses

Global

```
atomic_int x, y;  
atomic_init(&x, 1);  
atomic_init(&y, 1);
```

Thread 1

```
store_seq_cst(&x, 2);
```

Thread 2

```
store_seq_cst(&y, 2);
```

Threads 3 & 4

```
int a = load_seq_cst(&x);  
int b = load_seq_cst(&y);  
printf("%d %d", a, b);
```

Could print:

✓ 1 1

✓ 2 2

✓ **1 2**

✓ **2 1**

But **impossible to get both** 1 2 and 2 1.

The “happens before” relation

A **strict partial ordering** of events:

- Irreflexive: $\neg (A < A)$
- Asymmetric: $(A < B)$ implies $\neg (B < A)$
- Transitive: $(A < B \ \&\& \ B < C)$ implies $(A < C)$

In general, **no total ordering exists**, not even an unknown one. Some events are incomparable.

Axioms of happens-before

A happens before B when:

- A is **sequenced before** B in the same thread, according to ordinary language semantics, even if they're non-atomic.
- A has release ordering, B has acquire ordering, and the operations **synchronize with** each other. A synchronizes with B when the value read by B belongs to a **release sequence** headed by A.

Even if they're non-atomic? But wait, what about this?

Source

```
void foo(int *restrict a, int *restrict b) {  
    *a += 1;    *b += 1;    *a += 1;  
}
```

x86_64

```
inc    dword ptr [rsi]  
mov    eax, dword ptr [rdi]  
add    eax, 2  
mov    dword ptr [rdi], eax
```

aarch64

```
ldr    w8, [x1]  
ldr    w9, [x0]  
add    w8, w8, #1  
add    w9, w9, #2  
str    w8, [x1]  
str    w9, [x0]
```

riscv64

```
lw      a2, 0(a1)  
lw      a3, 0(a0)  
addiw   a2, a2, 1  
sw      a2, 0(a1)  
addiw   a3, a3, 2  
sw      a3, 0(a0)
```

Observability

- These reordering optimizations are legal only if they're *inobservable*.
- If no other thread is concurrently accessing $*a$ or $*b$, there's no observable change in program behavior.
- If some other thread *is* concurrently accessing them, that's a data race. Data races are UB, so the optimizer can ignore the possibility.

Release sequences

Global

```
atomic_int x;  
atomic_init(&x, 0);
```

Thread A

```
store_release(&x, 1);  
store_relaxed(&x, 2);  
fetch_add_relaxed(&x, 1);  
store_relaxed(&x, 6);
```

Thread B

```
fetch_add_release(&x, 1);  
store_relaxed(&x, 4);
```

Release sequences begin with a releasing store or RMW, continue with:

- A store by the releasing thread
- An RMW by any thread

Suppose our modification order is 1 2 3 4 5 6.
What are the release sequences?

Release sequences

Thread A

```
store_release(&x, 1);  
store_relaxed(&x, 2);  
fetch_add_relaxed(&x, 1);  
store_relaxed(&x, 6);
```

Thread B

```
fetch_add_release(&x, 1);  
store_relaxed(&x, 4);
```

Answer: 1 2 3 and 3 4 5.

- 2: same thread as 1, continues 1's sequence.
- 3: RMW continues 1's sequence, release starts a new one.
- 4: different thread from 1, breaks that sequence but continues 3.
- 5: RMW continues 3's sequence.
- 6: different thread from 3, breaks that sequence.

Release sequences begin with a releasing store or RMW, continue with:

- A store by the releasing thread
- An RMW by any thread

Suppose our modification order is 1 2 3 4 5 6.
What are the release sequences?

A simple spinlock

```
struct some_big_global_state foo;
atomic_flag foo_lock;

void lock(atomic_flag *lck) {
    while(!flag_test_and_set_acquire(lck));
}

void unlock(atomic_flag *lck) {
    flag_clear_release(lck);
}

void safely_frobnicate_foo() {
    lock(&foo_lock);
    frobnicate_nonatomically(&foo);
    unlock(&foo_lock);
}
```

Spinlock correctness

Thread 1..N

```
while(!flag_test_and_set_acquire(&foo_lock)); // 1
frob_nicate_nonatomically(&foo);                // 2
flag_clear_release(&foo_lock);                  // 3
```

- 1 Locking must happen before frobnicating.
- 2 Frobnicating must happen before unlocking.
- 3 Unlocking must happen before other threads lock.

- Reqs 1&2 satisfied by sequencing.
- Req 3: release on line 3 heads release sequence on &foo_lock; acquiring RMW on line 1 synchronizes w/ it, continuing only if it reads the 0 that a release wrote; extend by induction to any thread pair.

Fences: synchronizing without an address

Global

```
atomic_int x, y;  
atomic_init(&x, 0);  
atomic_init(&y, 0);
```

Thread 1

```
store_relaxed(&x, 1);    //1  
thread_fence_release(); //2  
store_relaxed(&y, 1);    //3
```

Thread 2

```
int a, b;  
b = load_relaxed(&y);    //4  
thread_fence_acquire(); //5  
a = load_relaxed(&x);    //6  
printf("%d %d", a, b);
```

Fences synchronize on *any* address written after a release fence and read before an acquire fence. These fences synchronize on *y*.

If $b = 1$ then $2 < 5$.

Fences: synchronizing without an address

Global

```
atomic_int x, y;  
atomic_init(&x, 0);  
atomic_init(&y, 0);
```

Thread 1

```
store_relaxed(&x, 1); //1  
thread_fence_release(); //2  
store_relaxed(&y, 1); //3
```

Thread 2

```
int a, b;  
b = load_relaxed(&y); //4  
thread_fence_acquire(); //5  
a = load_relaxed(&x); //6  
printf("%d %d", a, b);
```

Assuming $b = 1$:

- $2 < 5$
- By sequencing:
 - $1 < 2 < 3$
 - $4 < 5 < 6$
- By transitivity:
 $1 < 2 < 5 < 6$.
- Therefore, since
 $1 < 6, a = 1$.

We conclude that
printing 0 1 is
impossible.

Happens-before coheres with modification order

Four coherence properties:

- read/read
- read/write
- write/read
- write/write

If A happens before B , both act on a memory location M , A is a read (resp. write), and B is a read (resp. write), then the value read (resp. written) by A comes not later in M 's modification order than the value read (resp. written) by B .

mmap(2)

Map some virtual memory to

- Physical memory
- A file (on disk or on a tmpfs)
- Even device registers! (Linux uio facility)

Underlies

- `malloc(3)`
- `shm_open(3)` (semi-obsolete thanks to `mmap`)
- `brk(2)` / `sbrk(2)` (very obsolete)

Address translation example: RISC-V

- Page size: 2^{12} bytes = 4 KiB.
- Physical address: 56 bits = 64 PiB addressable = 2^{44} pages.
- Virtual address: $9 + 9 + 9 + 9 + 9 + 12 = 57$ bits, sign-extended to 64.
- Page table entry: 64 bits — 44-bit page number + 20 bits for flags or reserved. 2^9 entries fill one page.
- Top 9 bits of virtual address index into top-level page table, gives page number of next-level table.
- Repeat for each of 5 levels.
- 12-bit page offset untranslated (same in physical address as virtual).

Coping with SIGBUS

SIGBUS:

- Virtual address points to non-existent physical address, or past the end of a file
- Invalid address alignment
- Hardware error (bad sector, machine check exception)

(*c.f.* SIGSEGV: address not mapped at all, or without permissions for the operation)

In C:

- Call `setjmp()` before any access which might raise SIGBUS.
- Stash a pointer to the `jmp_buf` in thread-local storage.
- Recover by `longjmp()`ing out of the signal handler.
- Use `atomic_signal_fence()` as needed.

In Rust:

- Just use my `hw-exception` crate.

memfd_create(2)

- Returns an fd for an “anonymous file”.
- Pass the fd over a UNIX domain socket using SCM_RIGHTS (see `unix(7)`) to share it.
- Memfds can be “sealed” to prevent shrinking — no more worries about SIGBUS.

futex(2)

Spinlocks in userspace are rude! Use futex to schedule threads appropriately.

`futex(FUTEX_WAIT, addr, val, timeout):`

- Assert that `*addr` has value `val`.
- If not, return immediately.
- Otherwise, go to sleep for `timeout` or until another thread wakes us.

`futex(FUTEX_WAKE, addr, val):`

- Wake other threads, up to `val` of them, that are waiting on `addr`.

Lock-free algorithms

Standard terms:

- **Obstruction-free:** For all threads t , if every thread except t is suspended, t completes its operation in bounded time. (Livelock and starvation possible)
- **Lock-free:** There exists a thread which will complete its operation in bounded time, regardless of other suspensions. (Livelock impossible, starvation possible)
- **Wait-free:** All non-suspended threads will complete their operation in bounded time, regardless of other suspensions. (Livelock and starvation impossible)

My term: quasi-wait-free

- Assume distinct reader/writer roles.
- Suspended writers can still obstruct other threads...
- ...but no thread ever waits on a reader.

Sequence locks: a quasi-wait-free “blackboard”

```
struct seqlocked_state {  
    atomic_uint version;  
    struct state a_big_state_structure;  
}
```

One Writer

- 1 Increment version field even \rightarrow odd.
- 2 Update state.
- 3 Increment version field odd \rightarrow even.
- 4 FUTEX_WAKE on version field.

N Readers

- 1 Read version field.
- 2 Read state.
- 3 Read version field again.
- 4 If versions match and are odd, FUTEX_WAIT.
- 5 Loop until versions match and are even.

Sequence locks: a buggy attempt

Writer

- 1 Store-release odd version
- 2 Store-relaxed each word of state
- 3 Store-release even version

Reader

- 4 Load-acquire version
- 5 Load-relaxed each word of state
- 6 Load-acquire version

Problem: CPU can reorder steps 1/2, and 5/6.

- Suppose writer runs just once; version goes $0 \rightarrow 1 \rightarrow 2$.
- If step 4 reads `version=2`, then 3 happens before 4; we're good.
- But if steps 4&6 read `version=0`, no interthread happens-before relation is proven. Step 2 could be intermixed with 5.

Sequence locks: corrected

Writer

- 1 Store-relaxed odd version
- 2 Store-fence
- 3 Store-relaxed each word of state
- 4 Store-release even version

Reader

- 5 Load-acquire version
- 6 Load-relaxed each word of state
- 7 Load-fence
- 8 Load-relaxed version

- Store-fence prevents CPU reordering of 1/3; load-fence prevents reordering of 6/8.
- If any result of 3/4 is seen by 5/6, then 2 happens before 7, so 1 happens before 8.
- If 8 reads `version=0`, reason by modus tollens: 1 did not happen before 8, so 6 does not read anything written by 3.

A lock-free blackboard

```
struct blackboard {  
    atomic_uint version;  
    struct state a_big_state_structure[N]; // N>=2  
}
```

One Writer

- 1 Store-fence
- 2 Let $i = (\text{version} + 1) \% N$.
- 3 Store-relaxed i 'th copy of state.
- 4 Store-release increment to version field.

N Readers

- 1 Load-acquire version field.
- 2 Load-relaxed $(\text{version} \% N)$ 'th copy of state.
- 3 Load-fence
- 4 Load-relaxed version field again.
- 5 Loop until version increased by less than N.

Dealing with counter overflow

For k -bit version counter, we must assume fewer than 2^k writes occur while any single read is in progress. But in long-running systems, we must be robust against eventual overflow.

```
bool version_is_good(unsigned v1, unsigned v2) {  
    unsigned vmin = v1;  
    unsigned vmax = v1 + N; // OK to wrap  
  
    if (LIKELY(vmax > vmin)) {  
        //Normal case  
        return v2 >= vmin && v2 < vmax;  
    } else {  
        //Overflow case  
        return v2 >= vmin || v2 < vmax;  
    }  
}
```

Copying state: the safe, portable, slow way

In shared memory

```
struct shared_state {  
    atomic_uint version;  
    atomic_int field1;  
    atomic_char field2;  
    atomic_char field3;  
    //...  
}
```

In private memory

```
struct state {  
    unsigned int version;  
    int field1;  
    char field2;  
    char field3;  
    //...  
}
```

- Copy each field N one-by-one using `load_relaxed/store_relaxed`.
- Maximally portable, but will generate inefficient code.
- Annoying for library APIs: can't be generic over type of state.

Copying state: the fast, slightly dodgy way

```
void *racy_memcpy(void *restrict dest, const void *restrict src, size_t n) {  
    asm volatile(""); // Prevent optimizer from proving data race possibility  
    void *ret = memcpy(dest, src, n); // Full speed ahead, damn the torpedoes!  
    asm volatile("");  
    return ret;  
}
```

- `asm volatile` “black box” could stand in for lock acquisition/release, so no data race ever provable at compile time.
- Still UB in a real sense: architecture-dependent assumption that instructions used by `memcpy` don’t side-effect if a data race happens.

Q&A + further reading + contact

Primary sources:

- C17 standard: §5.1.2.4 (memory model) and §7.17 (<stdatomic.h>)
- Rust library reference: <https://doc.rust-lang.org/stable/std/sync/atomic>
- Your CPU's architecture manual
- Linux manpages: `mmap(2)`, `memfd_create(2)`, `futex(2)`, `unix(7)`

Secondary sources:

- C++ memory model:
https://en.cppreference.com/w/cpp/atomic/memory_order
- *Rust Atomics and Locks* by Mara Bos: <https://marabos.nl/atomics/>

These slides: <https://github.com/dfoxfranke/self2024>

Email: <dfoxfranke@gmail.com>

Appendix: What about `volatile`?

In theory:

- `volatile` has nothing to do with concurrency!
- Tells the compiler to treat load and stores as having side effects, so it won't reorder or combine the instructions.
- But the processor still can!
- Not atomic; does not prevent data races or UB.

In practice:

- Probably won't literally cause nasal demons.
- Adequately-aligned volatile pointers usually behave like relaxed atomics.
- On MSVC (only), volatile loads have acquire ordering and volatile stores have release ordering.
- But seriously, **just use real atomics**.

Appendix: Go sucks at concurrency!

- Channels: Go's favorite and only hammer
 - Real “share nothing” → never send a pointer over a channel → tons of unnecessary copying.
 - Sharing pointers → data races possible → not memory-safe any more. *E.g.* conflicting access to maps = heap corruption.
 - “CSP-inspired!” — Someone who's never read the CSP papers.
 - Looks more like π -calculus, but buffering and channel-closure semantics are kludgy.
- Go atomics: SeqCst is the only supported memory ordering. Order-of-magnitude performance penalty.
- Can escape into C and do low-level concurrency there, but Go FFI has high overhead: hundreds of nanoseconds per call into C. Have to do enough per call to make it worth it. Even Python doesn't have this problem!